

Kellin Pelrine

[Website](#) | [Google Scholar](#)

Leading cross-functional solutions on technical foundations.

Education

Ph.D. Machine Learning, McGill University, 2020-2025

Supervisors: Reihaneh Rabbany, Jean-François Godbout

M.A. Economics, Yale University, 2018-2019

M.S. Applied Mathematics, University of Colorado Boulder, 2017-2018

B.A. Economics and Mathematics, University of Colorado Boulder, 2014-2018

Class Rank 1, “Summa cum laude in economics” and “With Distinction”

Professional Experience

FAR AI

Member of Technical Staff, June 2025 – Present

Research Scientist, January 2025 – June 2025

Research Advisor, June 2023 – January 2025

Research Scientist Intern, December 2022 – June 2023

Leading the Integrity and Red teams, which aim to make AI trustworthy and secure.

Showed strongly superhuman capabilities will not be sufficient for robustness.

Paper 1 reported in [Financial Times](#), [The Times](#), [Ars Technica](#), [Vice](#)

Paper 2 reported in [Nature](#), [Ars Technica](#), [Scientific American](#)

Orals at ICML and AAAI, and [cited in US Senate testimony](#)

Called the “man who beat the machine”

Founded FAR.AI red-teaming function.

Our findings have influenced safeguards deployed in most of the leading AI systems, like ChatGPT, Claude, and Gemini.

Built partnerships with multiple frontier companies (e.g., OpenAI) and governments (e.g., the European Commission).

Built 0 → 5 people, now onboarded co-founder to continue scaling.

Leading projects to solve persuasion risks.

Built partnership with professors at MIT, Cornell, Mila, CMU, and CIDE.

Found Gemini would comply with requests to persuade people to join terrorist groups and other crimes – no jailbreaking needed – resulting in Google taking action to fix the gap in safeguards.

Discovered [jailbreak-tuning](#), the most severe blackbox fine-tuning attack.

Showed [increasingly capable AI is increasingly vulnerable](#) to data poisoning.

Secured \$4.1M in funding.

\$1.7M grant funding (900K PI, 800K co-PI)

\$2.4M contract funding (700K managing, 1.7M co-managing).

FAR.AI “Spirit” Award – for “energy, curiosity, and bold ambition”

Stitch

Cofounder and CTO, December 2023 – November 2024

Led development of a graphical interface for LLMs.

A visual system enables more organized, faster, and deeper interactions.

Mutual

Cofounder and Chief Scientist, June 2023 – December 2023

Led R&D to create consistent, reliable, and powerful generative AI agents.

Machine Learning Sports

Cofounder and CTO, October 2019 – December 2020

Created a recommendation system to help professional baseball pitchers optimize their pitch selection.

Go (strategy board game)

Professional teacher, freelance

Volunteer teacher, multiple Go clubs and online

Students have won tournaments and awards.

Other Research Leadership

Information Integrity in the GenAI Era

Project Director, January 2023 – August 2025

In coordination with PIs Reihaneh Rabbany and Jean-François Godbout, built tools to help people find reliable information and avoid being misled by manipulation from both humans and AI.

Initiated and supervised over a dozen subprojects: evidence retrieval, credibility assessment, simulations, uncertainty quantification, explainability, datasets, evaluation, fieldbuilding, and more.

Secured \$1.5M funding from Mila, Canadian Heritage Foundation, UKAIS, IVADO, FLI, CSDC, SPAR.

AI for Human Resilience

Project Director, February 2025 – August 2025

Demonstrated potential of AI as a fallible peer rather than authoritative teacher.

Science of building critical thinking, and AI and AI Safety literacy.

Workshop Organizer

Social Simulation with LLMs, COLM, [2025](#)

Temporal Graph Learning (TGL), NeurIPS, [2022](#) and [2023](#)

Opening remarks at first edition (2022).

These workshops led to the creation of a community with over 400 researchers, a recurring seminar series, and continuing series of workshops.

Axiom Futures Fellowship

Mentor, Summer 2024

Project (Simulations to Solve Societal-Scale Manipulation) went on to secure over \$500K funding, multiple publications, and an ongoing research agenda.

Mentees: Gayatri Krishna Kumar, Sneheel Sarangi.

Supervised Program for Alignment Research (SPAR)

Mentor, Spring+Fall 2023, Spring+Summer 2024, Spring+Fall 2025

Invited to mentor 40+ researchers in AI safety, from undergraduate through senior roles like Senior Applied Scientist, Applied Science Manager, and Staff SWE.

Mentees co-authored 8 published conference papers, dozens of workshop papers, and multiple works under review and in progress.

Mentees (ordered roughly chronologically): Caleb Gupta, Joel Christoph, Meilina Reksoprodjo, James Zhou, Lynn Feng, Mayank Goel, Raghav Ravi, Roman Hauksson, Tyler Vergho, Yury Orlovskiy, Arjun Verma, Arturs Semenuks, George Ingebretsen, Ruben Weijers, Gabrielle Castilho, Dylan Tabarini, Michael Walters, Will Cai, Ethan Kosak-Hine, Jason Zhang, Julius Broomfield, Reihaneh Iranmanesh, Sara Pieri, Tom Gibbs, Nikita Agarwal, Austin Welch, Toshali Goel, Kushal Dev, Luda Cohen, Sukanya Krishna, Hikaru

Tsujimura, Ardy Haroen, Deeraj Nagothu, Joshua Levy, Denis Volk, Anna Marchenkova, Igor Ivanov, Arth Singh, Mithil Srungarapu, Akash Kundu, Luis Ibanez, Adam Divak, Stephanie Ding, Yernat Yestekov.

Thesis Supervision

Co-supervisor, master's thesis of Ruben Weijers, Utrecht University, graduated 2025

Honors and Awards

Go (strategy board game)

U.S. Team Member, [World Mind Sports Games](#), 2012

95 countries competed. Invitation based on U.S. ranking and tournament results.

U.S. Team Member, [China-US Internet Go Tournament](#), 2020

One of 6 players selected by ranking to represent the U.S. in a friendship match with top Chinese amateurs.

Playoff for North American representative to Li Min Cup, 2014 and 2016

One of 8 participants by invitation.

Violin

Participant, 2 masterclasses of renowned pedagogue [Zakhar Bron](#), 2012

By invitation and audition.

Research and Academic (graduate)

Doctoral Training Scholarship, CAD 58k, Fonds de Recherche du Québec, 2023

Graduate Excellence Award, CAD 48k total, McGill University, 5 times, 2020-2023

GREAT Award, CAD 2k total, McGill University, 2 times, 2022 and 2023

IVADO PhD Excellence Scholarship, CAD 75k, IVADO, 2021

Max Stern Recruitment Fellowship, CAD 14k, McGill University, 2020

Cowles Foundation Fellowship, USD 32k, Yale University, 2018

Academic (undergraduate)

Chancellor's Recognition Award, CU Boulder, for perfect GPA, 2018

Sieglinde Talbott Haller Economics Scholarship, CU Boulder, 2016 and 2017

Jim and Laura Marshall Scholarship, CU Boulder, mathematics, 2016

Flock Leader Scholarship, CU Boulder, 2015

CU Esteemed Scholars Program: President Joseph A. Sewall Award, 2014

Selected Conference and Journal Papers *equal contribution, †equal advising

Open Technical Problems in Open-Weight AI Model Risk Management

Stephen Casper, Kyle O'Brien, Shayne Longpre, Elizabeth Seger, Kevin Klyman, Rishi Bommasani, Aniruddha Nrusimha, Ilia Shumailov, Sören Mindermann, Steven Basart, Frank Rudzicz, **Kellin Pelrine**, Avijit Ghosh, Andrew Strait, Robert Kirk, Dan Hendrycks, Peter Henderson, J Zico Kolter, Geoffrey Irving, Yarin Gal, Yoshua Bengio, Dylan Hadfield-Menell
TMLR (**Survey Certificate** for "exceptionally thorough or insightful survey")

Jailbreak-Tuning: Safeguards of Fine-Tunable Models are Illusory

Brendan Murphy, Dillon Bowen, Shahrhad Mohammedzadeh, Tom Tseng, Julius Broomfield, Adam Gleave, **Kellin Pelrine**
EMNLP 2025

A Guide to Misinformation Detection Datasets

Camille Thibault*, Jacob-Junqi Tian*, Gabrielle Péloquin-Skulski, Taylor Lynn Curtis, James Zhou, Florence Laflamme, Yuxiang Guan, Reihaneh Rabbany, Jean-François Godbout, **Kellin Pelrine**
KDD Datasets and Benchmarks Track 2025 (**Best Paper Runner Up**, top 3 of accepted papers)

Can Go AIs be adversarially robust?

Tom Tseng, Euan McLean, **Kellin Pelrine**[†], Tony Tong Wang[†], Adam Gleave[†]
AAAI 2025 (**Oral**)

Scaling Trends for Data Poisoning in LLMs

Dillon Bowen, Brendan Murphy, Will Cai, David Khachaturov, Adam Gleave[†], **Kellin Pelrine**[†]
AAAI 2025

Towards Reliable Misinformation Mitigation: Generalization, Uncertainty, and GPT-4

Kellin Pelrine, Anne Imouza, Camille Thibault, Meilina Reksoprodjo, Caleb Gupta, Joel Christoph, Jean-François Godbout, Reihaneh Rabbany
EMNLP 2023

Adversarial Policies Beat Superhuman Go AIs

Tony Tong Wang*, Adam Gleave*, Tom Tseng, **Kellin Pelrine**, Nora Belrose, Joseph Miller, Michael D Dennis, Yawen Duan, Viktor Pogrebniak, Sergey Levine, Stuart Russell
ICML 2023 (**Oral**)

Towards Better Evaluation for Dynamic Link Prediction

Farimah Poursafaei*, Andy Huang*, **Kellin Pelrine**, Reihaneh Rabbany
NeurIPS Datasets and Benchmarks Track 2022

The Surprising Performance of Simple Baselines for Misinformation Detection

Kellin Pelrine*, Jacob Danovitch*, Reihaneh Rabbany
The Web Conference 2021

Selected Workshop Papers

Emergent Persuasion: Will LLMs Persuade Without Being Prompted?

Vincent Chang, Thee Ho, Sunishchal Dev, Kevin Zhu, Shi Feng, **Kellin Pelrine**, Matthew Kowal
AIGOV @ AAAI 2026 (**oral**)

From Intuition to Understanding: Using AI Peers to Overcome Physics Misconceptions

Ruben Weijers, Denton Wu, Hannah Betts, Tamara Jacod, Yuxiang Guan, Vidya Sujaya, Kushal Dev, Toshali Goel, William Delooze, Reihaneh Rabbany, Ying Wu, Jean-François Godbout, **Kellin Pelrine**

ICLR 2025 Workshops: AI4CHL (**oral**), FM-WILD, Bi-Align

A Guide to Misinformation Detection Datasets

Camille Thibault*, Jacob-Junqi Tian*, Gabrielle Péloquin-Skulski, Taylor Lynn Curtis, James Zhou, Florence Laflamme, Yuxiang Guan, Reihaneh Rabbany, Jean-François Godbout, **Kellin Pelrine**

ICLR 2025 Workshops: MLDPR (**spotlight**), SCSL

A Simulation System Towards Solving Societal-Scale Manipulation

Maximilian Puelma Touzel*, Sneheel Sarangi*, Austin Welch*, Gayatri Krishnakumar, Dan Zhao, Zachary Yang, Hao Yu, Ethan Kosak-Hine, Tom Gibbs, Andreea Musulan, Camille Thibault, Busra Tugce Gurbuz, Reihaneh Rabbany, Jean-François Godbout, **Kellin Pelrine**

NeurIPS 2024 Workshops: SATA (**oral**), SoLaR, SafeGenAI

Selected Other Research

Large language models can effectively convince people to believe conspiracies

Thomas H. Costello, **Kellin Pelrine**, Matthew Kowal, Antonio A. Arechar, Jean-François Godbout, Adam Gleave, David Rand, Gordon Pennycook. 2026.

Securing Agentic AI - A Discussion Paper

Cyber Security Agency of Singapore, **FAR.AI**. 2025. (contributed as FAR.AI lead on this project)

It's the Thought that Counts: Evaluating the Attempts of Frontier LLMs to Persuade on Harmful Topics

Matthew Kowal, Jasper Timm, Jean-François Godbout, Thomas Costello, Antonio A. Arechar, Gordon Pennycook, David Rand, Adam Gleave, **Kellin Pelrine**. 2025.

Exploiting novel GPT-4 APIs

Kellin Pelrine*, Mohammad Taufeeque*, Michał Zając, Euan McLean, Adam Gleave. 2024.

Open, Closed, or Small Language Models for Text Classification?

Hao Yu, Zachary Yang, **Kellin Pelrine**, Jean Francois Godbout, Reihaneh Rabbany. 2023.

A Note on the Unconditional Bias of the Nadaraya-Watson Regression Estimator

Kellin Pelrine. Supervisor: Carlos Martins-Filho. Undergraduate Honors Thesis, 2018.