

Kellin Pelrine

[Website](#) | [Google Scholar](#)

Leading cross-functional solutions on a foundation of technical research.

Education

Ph.D. Machine Learning, McGill University, 2020-2025 (expected)

Supervisors: Reihaneh Rabbany, Jean-François Godbout

M.A. Economics, Yale University, 2018-2019

M.S. Applied Mathematics, University of Colorado Boulder, 2017-2018

B.A. Economics and Mathematics, University of Colorado Boulder, 2014-2018

Class Rank 1, “Summa cum laude in economics” and “With Distinction”

Scientific Objectives and Experience

Goal: Make generative AI a transformatively positive force for society, instead of a potentially catastrophically negative one.

Selected Technical Areas: AI security, AI agents, evaluation, reliability, robustness.

Selected Applications: Persuasion, manipulation, education, political polarization, Go.

Professional Experience

FAR AI

Member of Technical Staff, June 2025 – Present

Research Scientist, January 2025 – June 2025

Research Advisor, June 2023 – January 2025

Research Scientist Intern, December 2022 – June 2023

Leading the Integrity team, which aims to make AI trustworthy and secure.

Combined AI and Go expertise to study vulnerabilities in superhuman Go-playing AI.

Showed strongly superhuman capabilities will not be sufficient for robustness

Paper 1 reported in [Financial Times](#), [The Times](#), [Ars Technica](#), [Vice](#)

Paper 2 reported in [Nature](#), [Ars Technica](#), [Scientific American](#)

Orals at ICML and AAAI, and [cited in US Senate hearings](#)

The “man who beat the machine”

Leading projects uncovering frontier model vulnerabilities.

6 confidential long-form reports delivered as first or last (managing) author.
Exposed multiple critical vulnerabilities and numerous partial vulnerabilities.
Discovered [jailbreak-tuning](#), the most severe blackbox fine-tuning attack.
Showed [increasingly capable AI is increasingly vulnerable](#) to fine-tuning attacks.

Findings have influenced every major frontier model company and multiple governments.

Leading projects studying AI persuasion.

Forged collaboration with MIT, Cornell, Mila, American University, Centro de Investigación y Docencia Económicas.

Found Gemini would comply with requests to persuade to join terrorist groups and other crimes – minimal or no jailbreaking needed – resulting in Google taking action to fix the gap in safeguards.

Stitch

Cofounder and CTO, December 2023 – November 2024

Led development of a graphical interface for LLMs.

A visual system enables more organized, faster, and deeper interactions.

Mutual

Cofounder and Chief Scientist, June 2023 – December 2023

Led R&D to create consistent, reliable, and powerful generative AI agents.

Machine Learning Sports

Cofounder and CTO, October 2019 – December 2020

Created a pitch recommendation system powered by machine learning to help professional baseball pitchers optimize their pitch selection.

Go (strategy board game)

Professional teacher, freelance

Volunteer teacher, multiple Go clubs and online

Multiple tournament and award-winning students.

Research Leadership

Information Integrity in the GenAI Era

Project Director, January 2023 – Present

In coordination with PIs Reihaneh Rabbany and Jean-François Godbout, building tools that will help every person find reliable information and avoid being misled by malicious manipulation from both humans and AI.

Initiated and supervised over a dozen subprojects: evidence retrieval, credibility assessment, simulations, uncertainty quantification, explainability, datasets, evaluation, fieldbuilding, and more.

Secured \$1.5M funding from Mila, Canadian Heritage Foundation, UKAISI, IVADO, FLI, CSDC, SPAR.

AI for Human Resilience

Project Director, February 2025 – Present

Demonstrated potential of AI as a fallible peer rather than authoritative teacher. Science of building critical thinking, and AI and AI Safety literacy.

Workshop Organizer

Social Simulation with LLMs, COLM, [2025](#)

Future of Information Integrity Research (FIIR), ICDM, [2025](#)

Temporal Graph Learning (TGL), NeurIPS, [2022](#) and [2023](#)

Opening remarks 2022.

Axiom Futures Fellowship

Mentor, Summer 2024

Invited to mentor two full-time, funded fellowship recipients.

Project (Simulations to Solve Societal-Scale Manipulation) went on to secure over \$500K funding, multiple publications, one mentee joining the lab as a graduate student, and more work in progress.

Mentees: Gayatri Krishna Kumar, Sneheel Sarangi.

Supervised Program for Alignment Research (SPAR)

Mentor, Spring+Fall 2023, Spring+Summer 2024, Spring 2025

Invited to mentor 34 researchers in AI safety, from undergraduate through senior applied scientist.

Mentees co-authored 7 published conference papers, dozens of workshop papers, and multiple works under review and in progress.

Mentees (ordered roughly chronologically): Caleb Gupta, Joel Christoph, Meilina Reksoprodjo, James Zhou, Lynn Feng, Mayank Goel, Raghav Ravi, Roman Hauksson, Tyler Vergho, Yury Orlovskiy, Arjun Verma, Arturs Semenuks, George Ingebretsen, Ruben Weijers, Gabrielle Castilho, Dylan Tabarini, Michael Walters, Will Cai, Ethan Kosak-Hine, George Ingebretsen, Jason Zhang, Julius Broomfield, Reihaneh Iranmanesh, Sara Pieri,

Tom Gibbs, Nikita Agarwal, Austin Welch, Toshali Goel, Kushal Dev, Luda Cohen, Sukanya Krishna, Hikaru Tsujimura, Ardy Haroen, Deeraj Nagothu.

Thesis Supervision

Co-supervisor, master's thesis of Ruben Weijers, Utrecht University, graduated 2025

Honors and Awards

Research and Academic (graduate)

Doctoral Training Scholarship, CAD 58k, Fonds de Recherche du Québec, 2023

Graduate Excellence Award, CAD 48k total, McGill University, 5 times, 2020-2023

GREAT Award, CAD 2k total, McGill University, 2 times, 2022 and 2023

IVADO PhD Excellence Scholarship, CAD 75k, IVADO, 2021

Max Stern Recruitment Fellowship, CAD 14k, McGill University, 2020

Cowles Foundation Fellowship, USD 32k, Yale University, 2018

Academic (undergraduate)

Chancellor's Recognition Award, CU Boulder, for perfect GPA, 2018

Sieglinde Talbott Haller Economics Scholarship, CU Boulder, 2016 and 2017

Jim and Laura Marshall Scholarship, CU Boulder, mathematics, 2016

Flock Leader Scholarship, CU Boulder, 2015

CU Esteemed Scholars Program: President Joseph A. Sewall Award, 2014

Go (strategy board game)

U.S. Team Member, [World Mind Sports Games](#), 2012

95 countries competed. Invitation based on U.S. ranking and tournament results.

U.S. Team Member, [China-US Internet Go Tournament](#), 2020

One of 6 players selected by ranking to represent the U.S. in a friendship match with top Chinese amateurs.

Playoff for North American representative to Li Min Cup, 2014 and 2016

One of 8 participants by invitation.

Violin

Participant, 2 masterclasses of renowned pedagogue [Zakhar Bron](#), 2012

By invitation and audition.

Conference Papers *equal contribution, †equal advising

A Guide to Misinformation Detection Datasets

Camille Thibault*, Jacob-Junqi Tian*, Gabrielle Péloquin-Skulski, Taylor Lynn Curtis, James Zhou, Florence Laflamme, Yuxiang Guan, Reihaneh Rabbany, Jean-François Godbout, **Kellin Pelrine**

KDD Datasets and Benchmarks Track 2025 (**Best Paper Runner Up**)

The Structural Safety Generalization Problem

Tom Gibbs*, Julius Broomfield*, George Ingebretsen*, Ethan Kosak-Hine*, Tia Nasir, Jason Zhang, Reihaneh Iranmanesh, Sara Pieri, Reihaneh Rabbany, **Kellin Pelrine**
Findings of ACL 2025

Simulating public discourse in digital societies by giving social media to multimodal AI agents
Maximilian Puelma Touzel*, Sneheel Sarangi*, Gayatri Krishnakumar*, Busra Tugce Gurbuz,
Austin Welch, Zachary Yang, Andreea Musulan, Hao Yu, Ethan Kosak-Hine, Tom Gibbs, Camille Thibault, Reihaneh Rabbany, Jean-François Godbout[†], Dan Zhao[†], **Kellin Pelrine**[†]
IJCAI Demo Track 2025

Veracity: An Open-Source AI Fact-Checking System

Taylor Lynn Curtis, Maximilian Puelma Touzel, William Garneau, Manon Gruaz, Mike Pinder, Li Wei Wang, Sukanya Krishna, Luda Cohen, Jean-François Godbout[†], Reihaneh Rabbany[†], **Kellin Pelrine**[†]
IJCAI Demo Track 2025

Can Go AIs be adversarially robust?

Tom Tseng, Euan McLean, **Kellin Pelrine**[†], Tony Tong Wang[†], Adam Gleave[†]
AAAI 2025 (Oral)

Scaling Trends for Data Poisoning in LLMs

Dillon Bowen, Brendan Murphy, Will Cai, David Khachaturov, Adam Gleave[†], **Kellin Pelrine**[†]
AAAI 2025

Web Retrieval Agents for Evidence-Based Misinformation Detection

Jacob-Junqi Tian, Hao Yu, Yury Orlovskiy, Tyler Vergho, Mauricio Rivera, Mayank Goel, Zachary Yang, Jean-François Godbout, Reihaneh Rabbany, **Kellin Pelrine**
COLM 2024

Towards Reliable Misinformation Mitigation: Generalization, Uncertainty, and GPT-4

Kellin Pelrine, Anne Imouza, Camille Thibault, Meilina Reksoprodjo, Caleb Gupta, Joel Christoph, Jean-François Godbout, Reihaneh Rabbany
EMNLP 2023

SWEET: Weakly Supervised Person Name Extraction for Fighting Human Trafficking

Javin Liu*, Vidya Sujaya*, Peter Yu*, Pratheeksha Nair, **Kellin Pelrine**, Reihaneh Rabbany
Findings of EMNLP 2023

Party Prediction for Twitter

Kellin Pelrine, Anne Imouza, Gabrielle Desrosiers-Brisebois, Sacha Lévy, Jacob-Junqi Tian, Zachary Yang, Aarash Feizi, Cécile Amadoro, André Blais, Jean-François Godbout, Reihaneh Rabbany
ICWSM 2024

Adversarial Policies Beat Superhuman Go AIs

Tony Tong Wang*, Adam Gleave*, Tom Tseng, **Kellin Pelrine**, Nora Belrose, Joseph Miller, Michael D Dennis, Yawen Duan, Viktor Pogrebniak, Sergey Levine, Stuart Russell
ICML 2023 (Oral)

Towards Better Evaluation for Dynamic Link Prediction
Farimah Poursafaei*, Andy Huang*, **Kellin Pelrine**, Reihaneh Rabbany
NeurIPS Datasets and Benchmarks Track 2022

Extracting Person Names from User Generated Text: Named-Entity Recognition for Combating Human Trafficking
Yifei Li, Pratheeksha Nair, **Kellin Pelrine**, Reihaneh Rabbany
Findings of ACL 2022

The Surprising Performance of Simple Baselines for Misinformation Detection
Kellin Pelrine*, Jacob Danovitch*, Reihaneh Rabbany
The Web Conference 2021

Conference and Journal Papers Under Review

Jailbreak-Tuning: Safeguards of Fine-Tunable Models are Illusory
Brendan Murphy, Dillon Bowen, Shahrad Mohammedzadeh, Julius Broomfield, Adam Gleave,
Kellin Pelrine
ACL ARR 2025

Accidental Misalignment: Fine-Tuning Language Models Induces Unexpected Vulnerability
Punya Syon Pandey, Samuel Simko, **Kellin Pelrine**, Zhijing Jin
ACL ARR 2025

It's the Thought that Counts: Evaluating the Attempts of Frontier LLMs to Persuade on Harmful Topics
Matthew Kowal, Jasper Timm, Jean-François Godbout, Thomas Costello, Antonio A. Arechar,
Gordon Pennycook, David Rand, Adam Gleave, **Kellin Pelrine**
NeurIPS Datasets and Benchmarks 2025

Blueprint: A Social Media User Dataset for LLM Persona Evaluation and Training
Aurélien Bück-Kaeffer, Je Qin Chooi, Dan Zhao, Maximilian Puelma Touzel, **Kellin Pelrine**,
Jean-François Godbout, Reihaneh Rabbany, Zachary Yang
NeurIPS Datasets and Benchmarks 2025

Online Influence Campaigns: Strategies and Vulnerabilities
Andreea Musulan, Veronica Xia, Ethan Kosak-Hine, Tom Gibbs, Vidya Sujaya, Reihaneh Rabbany,
Jean-François Godbout[†], **Kellin Pelrine**[†]
Big Data and Society

Workshop Papers

Simulating public discourse in digital societies by giving social media to multimodal AI agents

Maximilian Puelma Touzel*, Sneheel Sarangi*, Gayatri Krishnakumar*, Busra Tugce Gurbuz, Austin Welch, Zachary Yang, Andreea Musulan, Hao Yu, Ethan Kosak-Hine, Tom Gibbs, Camille Thibault, Reihaneh Rabbany, Jean-François Godbout[†], Dan Zhao[†], **Kellin Pelrine[†]**
NLP4PI @ ACL 2025

From Intuition to Understanding: Using AI Peers to Overcome Physics Misconceptions
Ruben Weijers, Denton Wu, Hannah Betts, Tamara Jacod, Yuxiang Guan, Vidya Sujaya, Kushal Dev, Toshali Goel, William Delooze, Reihaneh Rabbany, Ying Wu, Jean-François Godbout, **Kellin Pelrine**
ICLR 2025 Workshops: AI4CHL (oral), FM-WILD, Bi-Align

Rethinking Anti-Misinformation AI
Vidya Sujaya, **Kellin Pelrine**, Andreea Musulan, Reihaneh Rabbany
ICLR 2025 Workshops: Bi-Align, HAIC

A Guide to Misinformation Detection Datasets
Camille Thibault*, Jacob-Junqi Tian*, Gabrielle Péloquin-Skulski, Taylor Lynn Curtis, James Zhou, Florence Laflamme, Yuxiang Guan, Reihaneh Rabbany, Jean-François Godbout, **Kellin Pelrine**
ICLR 2025 Workshops: MLDPR (spotlight), SCSL

A Simulation System Towards Solving Societal-Scale Manipulation
Maximilian Puelma Touzel*, Sneheel Sarangi*, Austin Welch*, Gayatri Krishnakumar, Dan Zhao, Zachary Yang, Hao Yu, Ethan Kosak-Hine, Tom Gibbs, Andreea Musulan, Camille Thibault, Busra Tugce Gurbuz, Reihaneh Rabbany, Jean-François Godbout, **Kellin Pelrine**
NeurIPS 2024 Workshops: SATA (oral), SoLaR, SafeGenAI

Epistemic Integrity in Large Language Models
Bijean Ghafouri, Shahrad Mohammadzadeh, James Zhou, Pratheeeksha Nair, Jacob-Junqi Tian, Mayank Goel, Reihaneh Rabbany, Jean-François Godbout, **Kellin Pelrine**
SafeGenAI @ NeurIPS 2024

The Structural Safety Generalization Problem
Tom Gibbs, Julius Broomfield, George Ingebretsen, Ethan Kosak-Hine, Tia Nasir, Jason Zhang, Reihaneh Iranmanesh, Sara Pieri, Reihaneh Rabbany, **Kellin Pelrine**
SafeGenAI @ NeurIPS 2024

Decompose, Recompose, and Conquer: Multi-modal LLMs are Vulnerable to Compositional Adversarial Attacks in Multi-Image Queries
Julius Broomfield, George Ingebretsen, Reihaneh Iranmanesh, Sara Pieri, Ethan Kosak-Hine, Tom Gibbs, Reihaneh Rabbany, **Kellin Pelrine**
NeurIPS 2024 Workshops: RBFM, Red Teaming GenAI

Scaling Laws for Data Poisoning in LLMs
Dillon Bowen, Brendan Murphy, Will Cai, David Khachaturov, Adam Gleave[†], **Kellin Pelrine[†]**
Workshop on Data Contamination @ ACL 2024

Can Go AIs be adversarially robust?
Tom Tseng, Euan McLean, **Kellin Pelrine[†]**, Tony Tong Wang[†], Adam Gleave[†]
Next Generation of AI Safety @ ICML 2024

Web Retrieval Agents for Evidence-Based Misinformation Detection

Jacob-Junqi Tian, Hao Yu, Yury Orlovskiy, Tyler Vergho, Mauricio Rivera, Zachary Yang,

Jean-François Godbout, Reihaneh Rabbany, **Kellin Pelrine**

Workshop on Online Harms and Abuse @ NAACL 2024

An Evaluation of Language Models for Hyperpartisan Ideology Detection in Persian Twitter

Sahar Omidi Shayegan, Isar Nejadgholi, **Kellin Pelrine**, Hao Yu, Sacha Levy, Zachary Yang,

Jean-François Godbout, Reihaneh Rabbany

Workshop on Resources and Technologies for Indigenous, Endangered and Lesser-resourced Languages in Eurasia (EURALI) @ LREC-COLING 2024

Uncertainty Resolution in Misinformation Detection

Yury Orlovskiy, Camille Thibault, Anne Imouza, Jean-François Godbout, Reihaneh Rabbany,

Kellin Pelrine

UncertaiNLP Workshop @ EACL 2024

Combining Confidence Elicitation and Sample-based Methods for Uncertainty Quantification in Misinformation Mitigation

Mauricio Rivera, Jean-François Godbout, Reihaneh Rabbany, **Kellin Pelrine**

UncertaiNLP Workshop 2024 @ EACL 2024

Quantifying learning-style adaptation in effectiveness of LLM teaching

Ruben Weijers, Gabrielle Fidelis de Castilho, Jean-François Godbout, Reihaneh Rabbany, **Kellin Pelrine**

Personalization of Generative AI Workshop 2024 @ EACL 2024

Comparing GPT-4 and Open-Source Language Models in Misinformation Mitigation

Tyler Vergho, Jean-François Godbout, Reihaneh Rabbany, **Kellin Pelrine**

Responsible Language Models (ReLM) @ AAAI 2024

Better Bridges Between Model and Real World

Kellin Pelrine

Canadian AI Conference Graduate Student Symposium 2023

Active Keyword Selection to Track Evolving Topics on Twitter

Sacha Lévy, Farimah Poursafaei, **Kellin Pelrine**, Reihaneh Rabbany

Workshop on Utility-Driven Mining and Learning @ ICDM 2022

OPPVIS: Visualizing Online Partisan Polarization of COVID-19

Zachary Yang, Anne Imouza, **Kellin Pelrine**, Sacha Lévy, Jiewen Liu, Gabrielle

Desrosiers-Brisebois, Jean-François Godbout, André Blais, Reihaneh Rabbany

IEEE Visualization & Visual Analytics 2021

Online Partisan Polarization of COVID-19

Zachary Yang, Anne Imouza, **Kellin Pelrine**, Sacha Lévy, Jiewen Liu, Gabrielle

Desrosiers-Brisebois, Jean-François Godbout, André Blais, Reihaneh Rabbany

International Conference on Data Mining Workshops (ICDMW) 2021

ComplexDataLab at WNUT-2020 Task 2: Detecting Informative COVID-19 Tweets by Attending over Linked Documents

Kellin Pelrine, Jacob Danovitch, Albert Orozco Camacho, Reihaneh Rabbany
Workshop on Noisy User-generated Text (WNUT) @ EMNLP 2020

Other Research

Exploiting novel GPT-4 APIs

Kellin Pelrine*, Mohammad Taufeeque*, Michał Zajac, Euan McLean, Adam Gleave. 2024.

Open, Closed, or Small Language Models for Text Classification?

Hao Yu, Zachary Yang, **Kellin Pelrine**, Jean Francois Godbout, Reihaneh Rabbany. 2023.

A Note on the Unconditional Bias of the Nadaraya-Watson Regression Estimator

Kellin Pelrine. Supervisor: Carlos Martins-Filho. Undergraduate Honors Thesis, 2018.

Referee

International Journal of Human-Computer Interaction

IEEE Transactions on Information Forensics and Security

ACM Computing Surveys

International Journal of Human-Computer Interaction

ICLR 2025

ACL ARR 2024, 2025

The Web Conference, 2023

NeurIPS Temporal Graph Learning Workshop, 2022 and 2023

NeurIPS, Datasets and Benchmarks Track, 2022

Workshop on Noisy User-generated Text, 2020

Invited Talks

Fine-Tunable AI Facilitates Evil Twins

Presented at INHR Track II AI Dialogue

Misinformation Detection with Generative AI

Research by: Jacob-Junqi Tian, Hao Yu, Yury Orlovskiy, Tyler Vergho, Mauricio Rivera, Mayank Goel, Zachary Yang, Jean-François Godbout, Reihaneh Rabbany, Kellin Pelrine

Presented at American Political Science Association Annual Meeting, 2024

Can Go AIs be adversarially robust?

Research by: Tom Tseng, Euan McLean, Kellin Pelrine[†], Tony Tong Wang[†], Adam Gleave[†]

Co-presenter: Tom Tseng

Presented at Mila AI Safety Reading Group

Towards Reliable Misinformation Mitigation: Generalization, Uncertainty, and GPT-4

Research by: Kellin Pelrine, Anne Imouza, Camille Thibault, Meilina Reksoprodjo, Caleb Gupta, Joel Christoph, Jean-François Godbout, Reihaneh Rabbany

Presented at American Political Science Association Annual Meeting, 2023

Presented at FAR Labs, 2023

Adversarial Policies Beat Superhuman Go AIs

Research by: Tony Tong Wang*, Adam Gleave*, Tom Tseng, Nora Belrose, Kellin Pelrine, Joseph Miller, Michael D Dennis, Yawen Duan, Viktor Pogrebniak, Sergey Levine, Stuart Russell

Co-presenter: Tony Tong Wang

Presented at Cross Labs, 2023

Party Prediction for Twitter

Research by: Kellin Pelrine, Anne Imouza, Gabrielle Desrosiers-Brisebois, Sacha Lévy, Jacob-Junqi Tian, Zachary Yang, Aarash Feizi, Cécile Amadoro, André Blais, Jean-François Godbout, Reihaneh Rabbany

Presented at Université de Montréal Political Science, 2023

Presented at IVADO Digital October, 2022

Presented at American Political Science Association Annual Meeting, 2022

Social Graphs

Guest Lecture

Presented at COMP 599 - Network Science, McGill University, 2022

Presented at COMP 599 - Network Science, McGill University, 2021

Political Polarization on Social Media

Research by: Kellin Pelrine, Anne Imouza, Gabrielle Desrosiers-Brisebois, Zachary Yang, Sacha Lévy, Aarash Feizi, Jiewen Liu, André Blais, Jean-François Godbout, Reihaneh Rabbany

Presented at American Political Science Association Annual Meeting, 2021

Using Social Media Data to Measure Polarization

Guest Lecture

Presented at PLU6904A - Les États-Unis de Trump à Biden, Université de Montréal/CÉRIUM, 2021

Marginal GAN

Research by: Kellin Pelrine

Presented at CU Boulder Econometrics Workshop, November 2019

ShapeAttack: Genetic Algorithm for Shape-Constrained Adversarial Robustness Testing

Research by: Kellin Pelrine

Presented at CU Boulder Econometrics Workshop, October 2019